

Get IT right



Puppet Ain't No Voodoo



Bartosz Majsak, Thomas Hug



◆ Bartosz Majsak

- ◆ Java Developer by day
- ◆ Open source junkie by night (Arquillian core team member)
- ◆ Conference speaker by passion (Devoxx, Jazoon ...)




◆ Thomas Hug

- ◆ With Cambridge Technology Partners since 2002
- ◆ Java Developer, TTL, Solution Architect
- ◆ Apache Committer, OSS contributor and aficionado



Why we recommended Linux?



- ◆ Infrastructure as code

- ◆ Version Control
- ◆ Reproducible



- ◆ Configuration Management

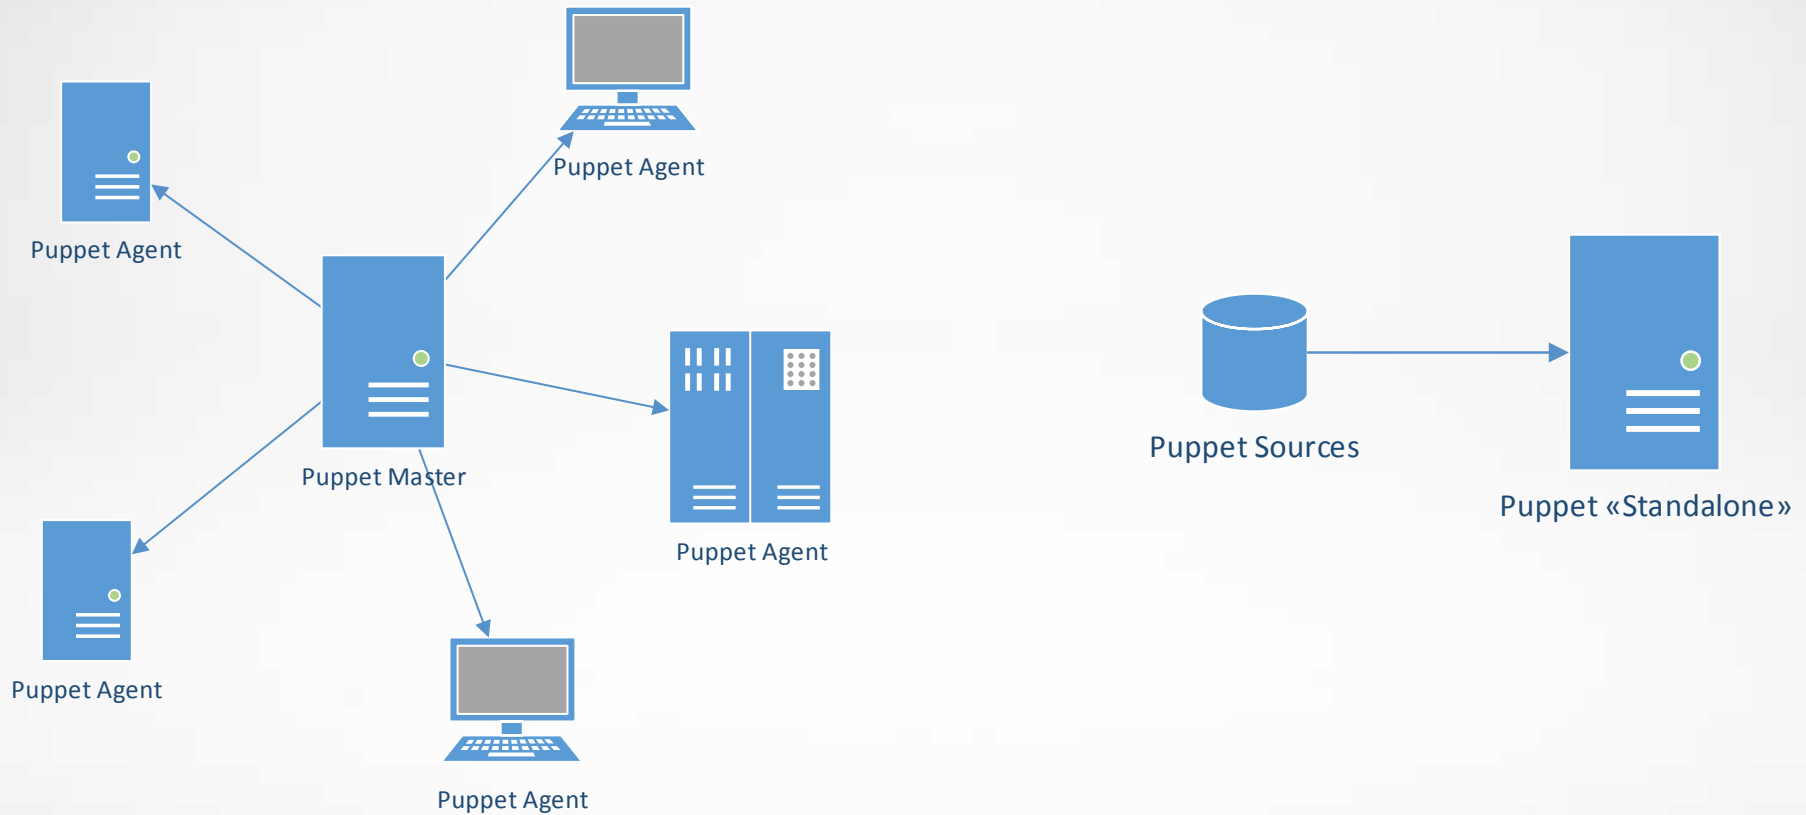
- ◆ Code = Documentation (if it's high-level enough)
- ◆ Readability
- ◆ Always in sync (if versioned and released properly)

- ◆ Automation

- ◆ Manual = slow and error prone
- ◆ Scalable over large installations



◆ Puppet Master/Agent vs. «Standalone» mode



- ◆ The Foreman
 - ◆ Advanced provisioning solution
 - ◆ ENC support
- ◆ Puppet Dashboards
 - ◆ Reporting
- ◆ Puppet Forge
 - ◆ Reusable Puppet modules
 - ◆ Share publicly

RED HAT®
SATELLITE



FOREMAN






- ◆ Official Puppet Documentation:
<https://docs.puppetlabs.com/>
- ◆ Puppet Type Reference:
<http://bit.ly/chopen-puppet-typeref>
- ◆ Puppet Cookbook:
<http://www.puppetcookbook.com/>
- ◆ You know this one anyway:
<http://stackoverflow.com/>



Getting Started



- ◆ Building a web server



◆ Ubuntu and Debian

```
wget http://apt.puppetlabs.com/puppetlabs-release-$(lsb_release -sc).deb  
sudo dpkg -i puppetlabs-release-$(lsb_release -sc).deb  
sudo apt-get update  
sudo apt-get -y install puppet ruby1.9.1-dev  
sudo gem install librarian-puppet puppet-lint
```

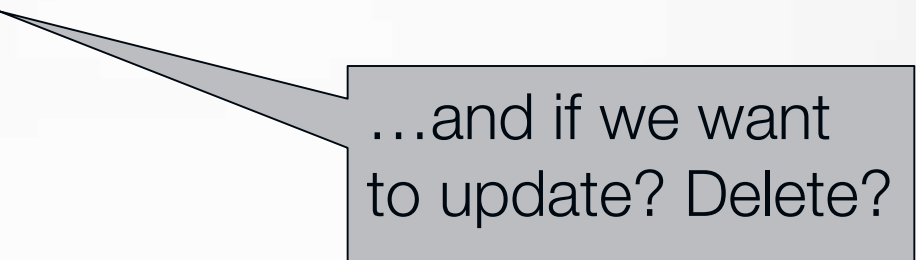
- ◆ Specific Ruby version is needed by Librarian

◆ Dedicated installers for Win and Mac OS X

- ◆ <http://downloads.puppetlabs.com/windows>
- ◆ <http://downloads.puppetlabs.com/mac/>

◆ The Shell Way:


```
if id -u 'tom' > /dev/null 2>&1; then
    echo "User exists, skipping"
else
    sudo adduser --disabled-password --gecos 'Tom Hug' tom
    echo "user created"
fi
```

A grey callout box with a black border and a pointer directed at the 'echo "user created"' line of the shell code. It contains the text "...and if we want to update? Delete?".

...and if we want to update? Delete?

◆ The Puppet Way:

```
user { 'tom':
    ensure    => exists,
    password => '*',
    comment  => 'Tom Hug'
}
```



- ◆ Find a list and reference at <https://docs.puppetlabs.com/references/latest/type.html>

◆ Setting up our site



```
$ subl site.pp  
  
file { '/tmp/hello.txt':  
    content => 'Hello world'  
}  
  
$ puppet apply site.pp  
...  
$ cat /tmp/hello.txt
```

◆ A typical resource:

```
RESOURCE { NAME:  
    ATTRIBUTE => VALUE,  
    ...  
}
```



◆ Interactive Labs

- ◆ We will get started with the solution...
- ◆ ...and leave some parts to complete
- ◆ Most build on top of each other!

◆ Full solution at GitHub

- ◆ <https://github.com/ctpconsulting/chopen-workshop-puppet-labs>
- ◆ Predefined checkpoints (tags) to advance / go back



```
$ git clone http://git.io/dX0Dxg labs_solution  
$ cd labs_solution  
$ git tag  
$ git checkout -f lab01
```




Part 1:

- ◆ Create a file `/etc/motd` and put in a quote using Puppet

Part 2 (use the type reference):

- ◆ Create a file `test` in `/tmp/puppet/files`
 - ◆ We expect your first attempt to fail ;-)
- ◆ Change the owner of the file to your current user
- ◆ Make it read-only

Part 3:

- ◆ Reorganize your `site.pp` file into folder structure `puppet/manifests`
- ◆ Rerun the script

Puppet Constructs

- ◆ Represents a host or host type



```
$ subl manifests/site.pp
```

```
node default {
```

```
  file { ['/tmp/hello.txt':  
        content => 'Hello World'  
      ]
```

```
}
```

```
$ cat /etc/hostname
```

- ◆ Can be inherited
- ◆ Name match with regular expressions

```
node 'xy' inherits 'z' { ... }
```

```
node /^(foo|bar)\.com$/ { ... }
```



Part 1:

- ◆ Change the nodename to match your host
- ◆ Verify Puppet runs your changes
 - ◆ Use the --debug flag to see what's going on

Part 2:

- ◆ Change the node name back to `default`
- ◆ Create another node with your hostname which inherits the default node
- ◆ Add an additional resource to verify inheritance

- ◆ Take advantage of the OS package manager



```
$ subl manifests/site.pp
```

```
node 'ubuntu' {
```

```
  package { 'nginx':  
    ensure => installed  
    #      ensure => 'major.minor.bugfix'  
    #      ensure => latest  
  }
```

```
  package { 'apache2.2-common':  
    ensure => absent  
  }
```

```
  package { 'nodejs':  
    ensure   => installed,  
    provider => 'npm'  
  }  
}
```

- Reusability?
- Scalability?

Other than the platform
default package manager

- ◆ Group of related resources



```
$ mkdir -p modules/nginx/manifest  
$ subl modules/nginx/manifests/init.pp
```

```
class nginx {  
    package { 'nginx':  
        ensure => installed  
    }  
}
```

```
class { 'nginx': }
```

```
$ sudo puppet apply modules/nginx/manifests/init.pp
```

- ◆ Group of related classes / definitions



```
$ subl modules/nginx/manifests/init.pp
```

```
class { 'nginx': }
```

```
$ subl manifests/site.pp
```

```
node 'ubuntu' {
```

```
    include nginx
```

```
}
```

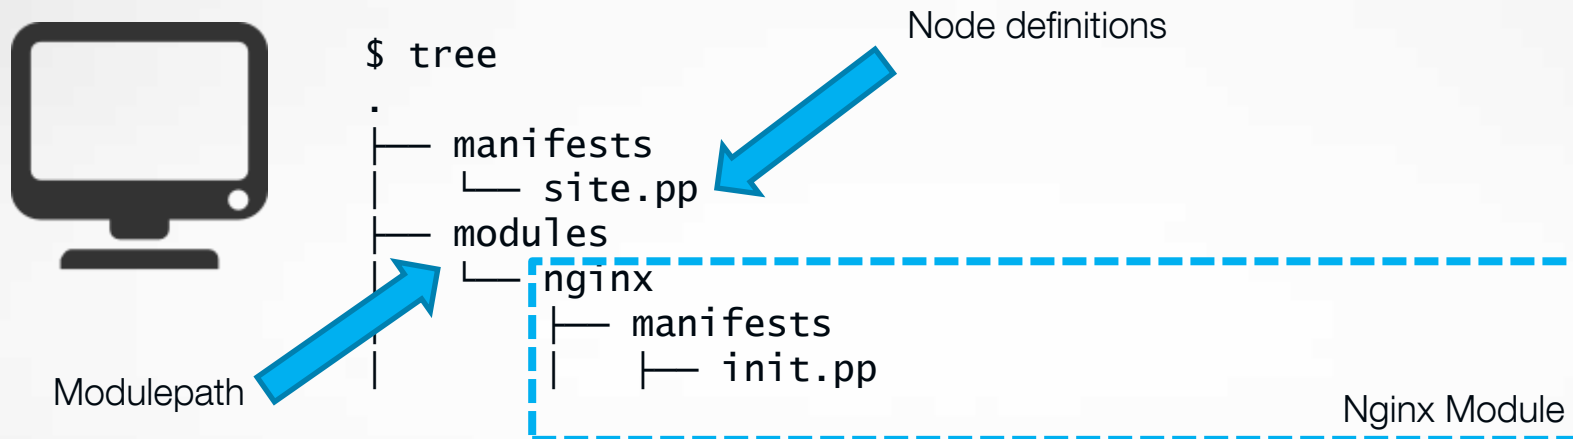
```
$ sudo puppet apply --modulepath=./modules manifests/site.pp
```

Autoload Format

- ◆ Modulepath: Where are my modules?

```
$ sudo puppet config print modulepath
```


- ◆ Node definition in separate site.pp
 - ◆ Execute this manifest
- ◆ Module encapsulating installation details
 - ◆ Make sure it's part of the module path





Part 1:

- ◆ Create the nginx module and main class
- ◆ Add the module to your node
- ◆ Make sure things still run as before

Part 2:

- ◆ Change the Puppet basemodulepath to find the module without using the command line param
- ◆ Find the Puppet config file and check what is in there
- ◆ You have now all the major Puppet building blocks!
 - ◆ Resources, Classes, Modules, Nodes



- ◆ Ensure the Nginx server is running



```
$ subl modules/nginx/manifests/init.pp
```

```
class nginx {
```

```
  package { 'nginx':  
    ensure => installed  
  }
```

```
  service { 'nginx':  
    ensure => running,  
    require => Package['nginx']  
  }
```

```
}
```

Resource / Class
dependencies

Dependencies can also be defined by:

- chaining resources with `-> ~>`
- declaring dependencies on collections:

```
Yumrepo<| |> -> Package<| provider == yum |>
```




- ◆ Run an arbitrary command



```
exec {  
  'set-licence-selected':  
    command => 'debconf-set-selections oracle-license accept true',  
    path     => ['/bin','/usr/bin'], user     => 'java',  
    unless  => 'debconf-get-selections | ...';  
  
  'set-licence-seen':  
    command => 'debconf-set-selections oracle-license seen true',  
    ...  
    creates => '/path/to/file';  
}
```

Insufficient **checks** on executions can make Puppet lose idempotence, or simply getting slow





Part 1:

- ◆ Add the nginx service to your manifests
 - ◆ Ensure it's running and started at boot time

Part 2:

- ◆ Refactor service and package into a separate class

Part 3:

- ◆ Run the script with the **--graph** option to see resource dependencies
 - ◆ Output folder is defined in the **graphdir** setting
 - ◆ Create a png with **dot -Tpng resource.dot > resource.png**

◆ Configure a site for nginx



```
$ mkdir -p modules/nginx/files  
$ subl modules/nginx/files/cat-pictures.conf
```

```
server {  
    listen 80;  
    root /var/www/cat-pictures;  
    server_name cat-pictures.com;  
}
```

Reusability?

```
$ subl modules/nginx/manifests/init.pp
```

```
class nginx {  
    ...  
    file { ['/etc/nginx/sites-enabled/default':  
        source => 'puppet:///modules/nginx/cat-pictures.conf',  
        notify => Service['nginx'],  
    }  
}
```

Puppet relative path
(no **files** in path)

Interested in changes



- ◆ Add the config and make sure all required paths exist
- ◆ Also add a simple index.html to the www root
- ◆ Access the server under <http://localhost>

◆ Reuse config files with templating



```
$ mv modules/nginx/files/cat-pictures.conf ../templates/vhost.conf.erb  
$ subl modules/nginx/templates/vhost.conf.erb
```

```
server {  
    listen 80;  
    root /var/www/<%= @site_name %>;  
    server_name /<%= @site_server %>;  
}
```

```
$ subl modules/nginx/manifests/init.pp
```

```
class nginx {  
  
    $site_name = 'cat-pictures'  
    $site_server = 'cat-pictures.com'  
  
    file { ['/etc/nginx/sites-enabled/default':  
        content => template('nginx/vhost.conf.erb'),  
        notify => Service['nginx'],  
    }  
}
```

Reusability?

◆ Reusable set of resources



Namespaces

```
$ subl modules/nginx/manifests/vhost.pp

define nginx::vhost($port, $site_name = $title) {
  $site_server = "${site_name}.com"

  file { ["/etc/nginx/sites-enabled/${site_name}":
    content => template('nginx/vhost.conf.erb'),
    notify => Service['nginx'],
  ]
}

$ subl modules/nginx/manifests/init.pp

class nginx {
  ...

  nginx::vhost { 'cat-pictures':
    port => 80
  }
}
```

Autoload format

Parameters
Default values for
optional parameters



Part 1:

- ◆ Change the configuration to a template
- ◆ Optional: Try using an **inline_template** without a file

Part 2:

- ◆ Apply the configuration in a definition
- ◆ Make sure the setup is working as expected

- ◆ After all it's a script



```
$ subl modules/nginx/manifests/vhost.pp

define nginx::vhost($site_name = $title, $port) {

    if ($port == undef or $port == 0) {
        fail('Port cannot be empty or 0!')
    }
    ...
}
```

- ◆ Use facts in classes, definitions or nodes

```
$ subl modules/nginx/manifests/init.pp
```

```
...
if ($::operatingsystem == 'windows') {
    notify { 'seriously?!': }
} else { ... }
```

Global Facts

```
$ facter | grep operating
```

◆ After all it's Ruby...



```
case $::operatingsystem {  
  'Debian', 'Ubuntu': {  
    ...  
  }  
  'windows': ...  
  default: ...  
}  
  
unless $::operatingsystem in 'windows' {  
  ...  
}  
  
if $::hostname =~ /app.*staging/ {  
  ...  
}  
  
$ec2_family = $::ec2_instance_type ? {  
  /t1/ => 'micro',  
  ...  
}
```



Part 1:

- ◆ Parameterize the nginx class with an optional version
- ◆ Use 'latest' if not set

Part 2:

- ◆ Install the nginx package from the Ubuntu PPA if running on Ubuntu (<http://wiki.nginx.org/Install>)
 - ◆ Lastet version is 1.6.0-1+trusty0
 - ◆ ppa:nginx/stable

- ◆ Clean unnecessary files with **tidy**



```
tidy { '/var/log/nginx':  
  age      => '4w',  
  size     => '100m',  
  recurse  => true,  
  matches  => [ "*.log", "*.tmp* ]  
}
```




- ◆ Clean out any existing files in the www directory before writing the actual files in there

Verifying Puppet

- ◆ Basic syntax checks with the **puppet parser** command



```
$ puppet parser validate --noop /path/to/manifest.pp

for file in $(find . -iname '*.pp'); do
    puppet parser validate --render-as s $file || exit 1;
done;
```

- ◆ Embeddable into
 - ◆ Jenkins / CI build
 - ◆ Commit hooks (e.g. Git)

- ◆ Avoid common mistakes, align formatting with **puppet-lint**



```
$ find . -iname *.pp -exec  
  puppet-lint --log-format  
    "%{path}:%{linenumber}:%{check}:%{KIND}:%{message}" {} \;
```

- ◆ Detailed usage instructions under <http://puppet-lint.com/>
- ◆ Follows style guide found under https://docs.puppetlabs.com/guides/style_guide.html

◆ Rspec-puppet and Cucumber-puppet

- ◆ BDD-style unit testing frameworks



```
$ puppet module generate chopen-rspectest
$ subl chopen-rspectest/spec/classes/init_spec.rb

require 'spec_helper'

describe 'rspectest' do
  context 'with defaults for all parameters' do
    it { should contain_class('rspectest') }
  end
end
```

- Rspec: <http://rspec-puppet.com/>
- Cucumber: <https://projects.puppetlabs.com/projects/cucumber-puppet/wiki>
- Articles on Puppet testing
 - <http://www.jedi.be/blog/2011/12/05/puppet-unit-testing-like-a-pro/>
 - <http://puppetlabs.com/blog/the-next-generation-of-puppet-module-testing>



Part 1:

- ◆ Run **puppet-lint** on your module
- ◆ Fix the errors and violations
 - ◆ Exclude the check for 80 chars per line

Part 2:

- ◆ Create a simple Rspec test
 - ◆ Copy template files from a temporary module
 - ◆ Check for a running nginx service
 - ◆ Run **rake help**



```
$ sudo apt-get install rake
$ sudo gem install rspec-puppet puppetlabs_spec_helper
...
$ rspec-puppet-init
$ echo "require 'puppetlabs_spec_helper/module_spec_helper'" >
spec/spec_helper.rb
```

Reinventing Wheels – NOT!

◆ DB to store cat pictures



```
$ puppet module generate ctp-mysql  
$ subl ctp-mysql/manifests/init.pp
```

```
class mysql {
```

```
  package { 'mysql-server':  
    ensure => installed,
```

```
  # Package -> File ~> Service pattern
```

```
  ... # umm, feels like someone should have done that already...?!
```

```
}
```


- ◆ Share and find some general usage Puppet modules:

<https://forge.puppetlabs.com/>



```
$ puppet module install puppetlabs-mysql
```



Part 1:

- ◆ Install the puppetlabs-mysql module
- ◆ Where is it installed to?
 - ◆ Make sure it is part of your modulepath

Part 2:

- ◆ Add a MySQL server to your nginx installation
- ◆ Add a DB and a user with access rights
 - ◆ Test with the mysql client

- ◆ Our Nginx module has now a dependency to MySQL
 - ◆ We need to document it
 - ◆ Or even better, ship the dependency declaration with the module



```
$ subl modules/nginx/metadata.json

{
  "name": "chopen-nginx",
  "version": "0.1.0",
  "author": "chopen",
  "summary": "Installs Nginx and the cat pictures website.",
  "dependencies": [
    {
      "name": "puppetlabs-mysql",
      "version_range": ">= 2.3.1"
    }
  ]
}
```

- ◆ No manual puppet module install
 - ◆ Endless and error prone
 - ◆ Transitive dependencies
- ◆ Librarian to the rescue: <http://librarian-puppet.com/>



```
$ subl Puppetfile
```

```
forge https://forgeapi.puppetlabs.com
```

```
def local(name)  
  mod "chopen/#{name}", :path => "./modules/#{name}"  
end
```

```
local 'nginx'
```

```
$ librarian-puppet install --path .librarian-modules  
$ sudo puppet apply manifests/site.pp --modulepath=.librarian-modules
```

- ◆ One primary repository
 - ◆ Contains site manifest with node definitions
 - ◆ Plus a Puppetfile for dependency management
- ◆ Modules outsourced to dedicated module repositories



```
$ subl Puppetfile

def git(name, version)
  mod "chopen/#{name}",
    :git => "git@server.com:#{name}.git",
    :ref => "#{version}"
end

git 'nginx', '0.1.0'
```





Part 1:

- ◆ Create a metadata.json file and fill in the dependency
 - ◆ You can get a skeleton file by running `puppet module generate`
- ◆ Create the Puppetfile and add a dependency to your Nginx module

Part 2:

- ◆ Add a Git dependency to the NTP module and add NTP to the node definition
 - ◆ Found at <https://github.com/puppetlabs/puppetlabs-ntp>

```
class { '::ntp':  
    servers => ['0.ch.pool.ntp.org', ... ],  
}
```

Environment Configuration

- ◆ MySQL user password
 - ◆ Different in production
 - ◆ Should preferably not show up in site.pp
- ◆ Parameterized class
 - ◆ Pattern: inherit param class where defaults are defined (e.g. OS-specific)



```
$ subl modules/nginx/manifests/init.pp

class nginx(

    $mysql_user = $nginx::params::mysql_user,
    $mysql_password = undef

) inherits nginx::params {

    ...

}
```

◆ Hiera = Hierarchical Database

- ◆ Integrated with Puppet
- ◆ Can have various backends, from file-based to DBs to REST endpoints
- ◆ Externalized: Can be managed separately



```
$ mkdir -p hiera/node  
$ subl hiera/hiera.yaml
```

```
---  
:backends:  
  - yaml  
:yaml:  
  :datadir: hiera  
:hierarchy:  
  - "node/%{::hostname}"
```

```
$ sudo puppet config print hiera_config  
$ subl hiera/node/ubuntu.yaml
```

```
---  
nginx::mysql_password: test
```

```
$ hiera -c hiera/hiera.yaml nginx::mysql_password ::hostname=ubuntu
```



Part 1:

- ◆ Create a hiera.yaml file with a yaml backend
- ◆ Add a hierarchy containing your node settings
- ◆ Configure Puppet to use the new Hiera config

Part 2:

- ◆ Optional: Use a MySQL backend
 - ◆ Found at <https://github.com/Telmo/hiera-mysql-backend>

Advanced Topics



- ◆ Config Files might be part of a package
 - ◆ Changing the template with each new version? Nope.
 - ◆ String replacement in an XML file? Nope.
 - ◆ Augeas adds advanced editing capabilities



```
<Server port="8080" schema="http"> ...

augeas { "tomcat-settings":
  lens      => "Xml.lns",
  incl      => $tomcat::configfile,
  changes   => [
    "set Server/#attribute/port ${tomcat::port}",
    "set Server/#attribute/schema https",
    "set Server/#attribute/proxyName server.chopen.ch",
    "set Server/#attribute/proxyPort 443",
  ],
  require   => File["${tomcat::configfile}"],
}
```

- ◆ Ship reusable facts with a module
 - ◆ Requires some custom Ruby code though




```
$ mkdir modules/nginx/lib/facter
$ subl modules/nginx/lib/facter/user_home.rb
```

Autoload format

```
# A facter fact to determine the user home directory.
Facter.add("user_home") do
  setcode do
    Facter::Util::Resolution.exec("echo $HOME")
  end
end
```

- ◆ Come back to next year's workshop 😊
 - ◆ Out of scope, but good to know you can
 - ◆ O'Reilly book reference, or GitHub



- ◆ When resource dependencies get out of control...
 - ◆ Assign classes to custom stages (default is «main»)
 - ◆ Stages freely orderable
 - ◆ Use carefully!



```
class stages {  
  
  stage { 'first': before => Stage['main'] } # consider site.pp  
  stage { 'last': require => Stage['main'] } # to define stages  
  
  class me_first { notify { 'In first!': } }  
  class me_last { notify { 'Out last...': } }  
  
  class { 'me_first':  
    stage => 'first',  
  }  
  class { 'me_last':  
    stage => 'last',  
  }  
  
}
```


Get IT right

Thank you!



- ◆ Icons provided by Icons8: <http://icons8.com/>